



DECKHOUSE

**Kubernetes
Platform**

L2 Load Balancer Basics

L2LoadBalancer



There are three frontends nodes and one worker in the cluster.

L2LoadBalancer



```
apiVersion: deckhouse.io/v1alpha1
kind: ModuleConfig
metadata:
  name: l2-load-balancer
spec:
  enabled: true
  version: 1
```



We turn on l2-load-balancer module.

L2LoadBalancer



```
apiVersion: deckhouse.io/v1alpha1
kind: L2LoadBalancer
metadata:
  name: front
spec:
  addressPool:
    - 192.168.122.100-192.168.122.150
  nodeSelector:
    node-role: front
```



An L2LoadBalancer resource has been created specifying front-end nodes and a pool of "public" IP addresses. This allows for easily creating "zones" by associating specific address pools with a group of nodes. Speakers are run on all front-end nodes.

L2LoadBalancer




```
apiVersion: v1
kind: Service
metadata:
  name: nginx-deployment
  annotations:
    network.deckhouse.io/l2-load-balancer-name: ingress
    network.deckhouse.io/l2-load-balancer-external-ips-count: "3"
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: nginx
  type: LoadBalancer
```



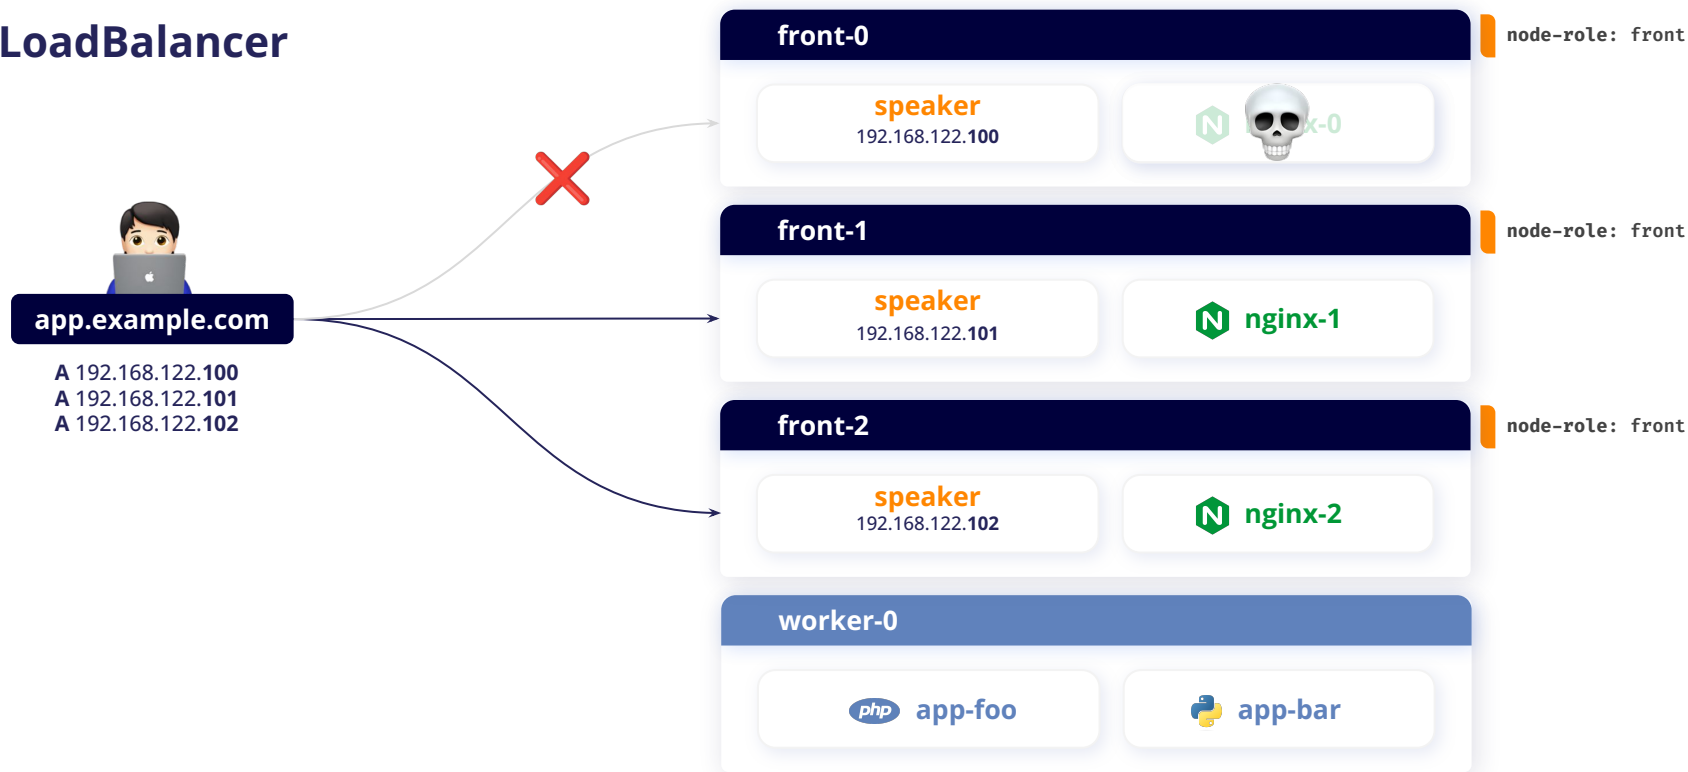
An **Service** resource with type **LoadBalancer** has been created. It contains special annotations with the name of the L2LoadBalancer and the required number of IP addresses. Speakers are launched on all frontend nodes, each obtaining one or more addresses from the pool.

L2LoadBalancer



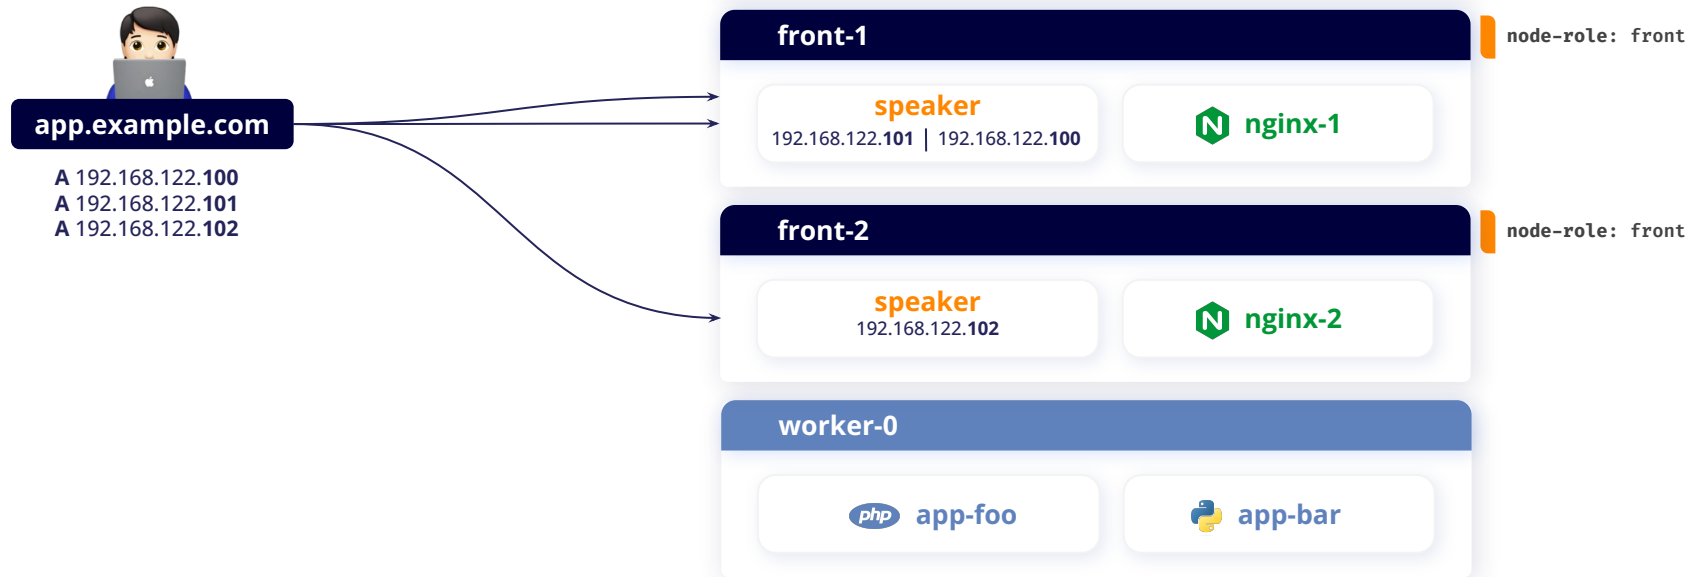
Each front-end node participates in handling application requests.
For this, three A records are specified in the public DNS name of the application.

L2LoadBalancer



In the event of a failure of the nginx application on one of the front-end nodes or the node itself, a third of the requests will fail,...

L2LoadBalancer



...and one of the remaining front-end nodes will take over the "problematic" IP address and handle the incoming application requests.